

SRP #11: Secure Remote Passwords behind PKCS #11

Rick van Rein
OpenFortress

December 20, 2015

Abstract

Secure Remote Passwords provide a zero-knowledge proof system for password-protected service access; furthermore, the protocol establishes a session key with forward secrecy properties. These properties make it a desirable alternative to many other password schemes; its integration as a TLS cipher suite also makes it practically usable. In this paper, we find what is required to use SRP with a randomly generated password that is protected by a PKCS #11 token.

1 Introduction

The SRP formalism is related to Diffie-Hellman, but it is not the same. This means that implementation on a PKCS #11 token, with its stringent assumptions on the use of DH and the accompanying protection of secrets, is not trivial. In fact, as will be shown here, we need to modify the client side of the protocol to be able to use SRP. Since only the client-side changes and the cryptographic properties of the protocol are not impacted, this is a reasonable choice to consider.

Formalism of SRP. The following exchange is defined for SRP. Client and service are assumed to agree on DH parameters, a generator g and a hash function H . All equalities in this paper are modulo the DH modulus. The password is fixated by the client as P , which may be shared among many services without loss of security. For each service, the client generates a random salt s and passes it with a verifier v to the service while setting up an account C . So, the service stores $C \mapsto \langle s, v \rangle$ and treats v as a secret, because it might elicit password cracking attempts. The client stores C , which it sends to the service to get back the salt s that it setup in the past.

Client	Protocol	Service
--------	----------	---------

$$\begin{aligned}
x &= H(s, P) \\
v &= g^x \\
C, s, v &\longrightarrow \text{store } C \mapsto \langle s, v \rangle
\end{aligned}$$

During authentication, the client picks random value a and the service picks random values u and b . We assume the entropy of each of these three values to match the size of the to-be-derived session key. The authentication protocol is conducted as follows:

Client	Protocol	Service
	$C \longrightarrow$	lookup $\langle s, v \rangle$
	$\longleftarrow s$	
$x = H(s, P)$		
$A = g^a$		
	$A \longrightarrow$	
$k = H(N, g)$		$k = H(N, g)$
		$B = k \cdot v + g^b$
	$\longleftarrow B$	
$u = H(A, B)$		$u = H(A, B)$
$S = (B - k \cdot g^x)^{a+ux}$		$S = (A \cdot v^u)b$
$K = H(S)$		$K = H(S)$
$M_1 = H(H(N) \oplus H(g), H(C), s, A, B, K)$		
	$M_1 \longrightarrow$	verify M_1
		$M_2 = H(A, M_1, K)$
	verify M_2 $\longleftarrow M_2$	

Rules of the game. We define the following constraints to the work laid out in this paper:

- The purpose of using PKCS #11 is to protect credentials from extraction, so that the SRP functions of the client cannot be implemented without the PKCS #11 token holding the password.
- The random value a is not protected by PKCS #11 because it only impacts the session key, and this value is also known to the service.
- Intermediate values will be analysed; they are considered ‘public’, that is need no PKCS #11 protection, when they can be derived on the service-side of the protocol. Note however, that a passive observer sees less.

- We only define client-side computations with PKCS #11, because the the service side does not handle any secrets, thanks to the zero-proof nature of the protocol.
- We follow the assumption of SRP that the verifier is kept a secret by the service side; this is chiefly in the interest of the service, the client could employ SRP salt pinning to enforce it. Salt pinning falls outside the scope of this paper.

2 First Naive Implementation

To demonstrate the problems that arise when implementing unchanged SRP on a PKCS #11 interface, we start off by making a sound technical implementation, and point out where this fails to satisfy the desired password protection for which PKCS #11 is being considered.

Password establishment. A new password is generated with the DH mechanism in `C.GenerateKeyPair()` in such a way that it must remain on the token. This will be the password P for the formalism. Since the client performing this computation is also the party to generate salt and verifier, and to authenticate on later returning, there should be no conflict of interest to support key export. Only for reasons of backup might wrapped export be permissible.

Credentials generation. To derive service access credentials from the password, first generate a random salt s . The value is public, and may be generated on or off the PKCS #11 token. Then, the token's `C.DigestXXX()` functions must be used to compute the hash value $x = H(s, P)$ on the token; the password can be incorporated with `C.DigestKey()` pointing at P .

Now run `C.DeriveKey()` with g as the base value and point to x as the private key or exponent. The outcome is $v = g^x$ which is supplied in a pair $\langle s, v \rangle$ to the service, together with a username.

Authentication. The client starts off as basic DH with `C.GenerateKeyPair()`. The private key will be known as a and the public key is A . Send A over to the service.

The service generates $B = v + g^b$ for a local random secret b and sends B and u to the client, where u is another random scrambling parameter that is publicly shown.

The client now computes $B - g^x = B - v$, which is entirely based on public values. The result equals g^b , which can be combined with session key a in `C.DeriveKey()` to find g^{ab} , a normal DH result.

Then, B^u is computed from public values, and used as a base key in `C.DeriveKey()` with DH private key x . This yields $B^{ux} = g^{bux}$, a value that can also be calculated separately on the service, as $v^{bu} = g^{xbu}$.

The values g^{ab} and g^{bux} are now multiplied outside of PKCS #11 to find shared secret S , which is of course also derived on the service. Likewise, the session key $K = H(S)$ is also available to the service and not considered a secret that would benefit from PKCS #11 protection.

Problem Analysis. The value of this implementation is that it never needs to extract the password P , so PKCS #11 can be exploited as a barrier to the secret. However, this implementation calculates x from a `C.DigestFinal()` invocation, which means it is delivered into a byte array that is extracted from PKCS #11. The value x is the only part of the protocol that contains P , so knowing x is as useful as knowing P , at least for a service that uses the given salt value s .

In fact, in terms of the original AKA formalism of which SRP is an instance, the computation of x falls outside the formalism as a mere parameter. The use of x in the SRP protocol may be considered a technical aspect, rather than a cryptographic necessity. This means that we may consider modifications to the calculation of x to make it better suited for PKCS #11 implementation.

3 Client-side Variation on the Formalism

The proposed variation to better adapt SRP to PKCS #11 is to use another one-way function than $H(s, P)$. Within the SRP formalism, the operations for the DH mechanism offers an alternative, namely the PKCS #11 function `C.DigestKey()` which can be used as a modular-exponentiation operation, possibly incorporating secret exponentials.

As an alternative to x in plain SRP, we define $x' = H(s)P$ to improve the integration. This looks awful, because it is possible to derive P from the values of x' and s , but the damage is controlled by not actually depending on the value x' in the implementation; instead, we use it to modify the formalism. Note that the hash function has no cryptographic role in the AKA formalism, so we might even have specified sP instead of $H(s)P$. We choose to retain the hash because the salt may be used to carry additional values to pin down the service, and that could be computationally expensive or run into size constraints without the hash function.

During secret establishment, the client does not generate a random P on

the token, but instead constructs a DH key pair $\langle p, P \rangle$:

$$p = g^P$$

Both P and p are concealed by PKCS #11, with no permission of extracting them. Wrapped export might be supported to facilitate backup and recovery, but the values of p and P are assumed to not appear in plaintext outside of PKCS #11.

When setting up a service with a verifier and salt, the following process is used:

Client	Protocol	Service
$x' = H(s)P$		
$v' = p^{H(s)}$		
	$C, s, v' \longrightarrow$	store $C \mapsto \langle s, v' \rangle$

The service cannot tell the difference, because the values s is treated as an opaque string and because it is not of a different format; the value v' is also not distinguishable for the service.

During authentication, the client picks random numbers a and the service picks random values u and b . The authentication protocol is conducted as follows:

Client	Protocol	Service
	$C \longrightarrow$	lookup $\langle s, v' \rangle$
	$\longleftarrow s$	
$v' = p^{H(s)}$		
$A = g^a$		
	$A \longrightarrow$	
$k = H(N, g)$		$k = H(N, g)$
		$B = k \cdot v' + g^b$
	$\longleftarrow B$	
$u = H(A, B)$		$u = H(A, B)$
$S' = (B - k \cdot v')^a \cdot (((B - v')^u)^{H(s)})^P$		$S = (A \cdot v'^u)^b$
$K' = H(S')$		$K = H(S)$
$M_1 = H(H(N) \oplus H(g), H(C), s, A, B, K)$		
	$M'_1 \longrightarrow$	verify M'_1
		$M'_2 = H(A, M'_1, K)$
verify M'_2	$\longleftarrow M'_2$	

The calculation procedure for S' by the client differs from the original calculation, but ought to yield the same output as the standard calculation of S by the service under proper authentication conditions. The same then holds for K' , M'_1 and M'_2 . We only use the accents to the symbols to distinguish their calculation and be able to prove equality.

Note that the value x' is never seen in this calculation, nor can it be derived. The value P is used, but cannot be extracted from PKCS #11.

In terms of the original work on SRP, this is still an AKA formalism. The change from $x = H(s, P)$ to $x' = H(s)P$ occurs in the periphery surrounding the AKA formalism definitions. We will consider efficiency and security impacts in subsequent sections, but first turn to an implementation in terms of PKCS #11, using standard DH mechanisms.

Proof of correctness. During authentication, the two sides need to find the same values for S and S' , after which the values for K' , M'_1 and M'_2 follow without a change. The thing to prove is that S' in the client-side formalism yields the same result as S in the service-side formalism, given that the service has received the value v' instead of v . The correctness is proven as follows:

$$\begin{aligned}
S' &= (B - k \cdot v')^a \cdot (((B - k \cdot v')^u)^{H(s)})^P \\
&= (g^b)^a \cdot (g^b)^{uH(s)P} \\
&= g^{ba} \cdot g^{buH(s)P} \\
&= g^{ab} \cdot ((g^P)^{H(s)})^{ub} \\
&= (g^a)^b \cdot (p^{H(s)})^{ub} \\
&= A^b \cdot v'^{ub} \\
&= (A \cdot v'^u)^b \\
&= S
\end{aligned}$$

This establishes that the unmodified SRP service algorithm, as well as the unmodified networking protocol, can work with the modified client that this formalism introduces. The one thing to care for is that all clients adhere to the same new formalism, since they will need to base their work on v' and not v . Since the password pair $\langle p, P \rangle$ is concealed by PKCS #11, this constraint is already implied by the technology.

Problem Analysis. The formalism presented here suffers from a cryptographic problem, namely that the value p is assumed to be secret. In normal Diffie-Hellman operations, only P would be considered secret, and p would be a public key. Indeed, the value of p could be recomputed from g and P by anyone with PKCS #11 access, and extracted. The impact of this extraction would be that new values of v' could be calculated without the help of PKCS #11; only authentication would continue to rely on live access to

PKCS #11, but that is not the best possible result. We therefore construct one further iteration.

Another problem that remains is the potential relationship between values of v . Consider two values of $H(s)$, called h_1 and h_2 . Since these are random values, it may happen that h_2 is divisible by h_1 , so when $v_1 = p^{h_1}$ then $v_2 = p^{h_2} = v_1^{h_2/h_1}$. This problem is mitigated in plain SRP through the use of the hash function, and the next iteration must also resolve it through better scattering of the verifiers.

4 Maximally PKCS #11 Protected Formalism

On top of the suitability for PKCS #11 in previous variation, the additional requirement is now to avoid that verifiers can be constructed from extractable values like p . This is ensured by applying the secret key P to the salt.

As an alternative to x and x' , we therefore define $x'' = H(s)^P P$. This introduces a multiplication with secret P , which will make it impossible to derive password-equivalent value x'' in the implementation. It also raises the salt to the power P to require the password before the verifier can be computed. Under the assumption that services conceal the verifier, a new service cannot be setup with the same salt, and the dependency on P implies that only the client with PKCS #11 access can thus create a new salt and verifier for a new service.

The secret key P established as before, with token-generated random key P , protected from extraction. But in this variation, the value p does not have to be protected from extraction.

$$p = g^P$$

When setting up a service with a verifier and salt, the following process is used:

Client	Protocol	Service
$x'' = H(s)^P P$		
$v'' = p^{H(s)^P}$		
	$C, s, v'' \longrightarrow$	store $C \mapsto \langle s, v'' \rangle$

As before, the service cannot tell the difference, because the value s is treated as an opaque string and because it is not of a different format; the value v'' is also not distinguishable for the service.

During authentication, the client picks random value a and the service picks random values u and b . The authentication protocol is conducted as follows:

Client	Protocol	Service
	$C \longrightarrow$	lookup $\langle s, v'' \rangle$
	$\longleftarrow s$	
$v'' = p^{H(s)^P}$		
$A = g^a$		
	$A \longrightarrow$	
$k = H(N, g)$		$k = H(N, g)$
		$B = k \cdot v'' + g^b$
	$\longleftarrow B$	
$u = H(A, B)$		$u = H(A, B)$
$S'' = (B - k \cdot v'')^a \cdot (((B - k \cdot v'')^u)^{H(s)^P})^P$		$S = (A \cdot v''^u)^b$
$K'' = H(S'')$		$K = H(S)$
$M_1 = H(H(N) \oplus H(g), H(C), s, A, B, K)$		
	$M_1'' \longrightarrow$	verify M_1''
		$M_2'' = H(A, M_1'', K)$
	verify M_2'' $\longleftarrow M_2''$	

The calculation procedure for S'' by the client differs from the original calculation, but ought to yield the same output as the standard calculation of S by the service under proper authentication conditions. The same then holds for K'' , M_1'' and M_2'' . We only use the accents to the symbols to distinguish their calculation and be able to prove equality.

Note that the value x'' is never seen in this calculation, nor can it be derived. The value P is used, but cannot be extracted from PKCS #11.

In terms of the original work on SRP, this is still an AKA formalism. The change from $x = H(s, P)$ to $x'' = H(s)^P P$ occurs in the periphery surrounding the AKA formalism definitions. We will consider efficiency and security impacts in subsequent sections, but first turn to an implementation in terms of PKCS #11, using standard DH mechanisms.

Proof of correctness. During authentication, the two sides need to find the same values for S and S'' , after which the values for K'' , M_1'' and M_2'' follow without a change. The thing to prove is that S'' in the client-side formalism yields the same result as S in the service-side formalism, given that the service has received the value v'' instead of v . The correctness is proven as follows:

$$S'' = (B - k \cdot v'')^a \cdot (((B - k \cdot v'')^u)^{H(s)^P})^P$$

$$\begin{aligned}
&= (g^b)^a \cdot (g^b)^{uH(s)^P P} \\
&= g^{ba} \cdot g^{buH(s)^P P} \\
&= g^{ab} \cdot ((g^P)^{H(s)^P})^{ub} \\
&= (g^a)^b \cdot (p^{H(s)^P})^{ub} \\
&= A^b \cdot v''^{ub} \\
&= (A \cdot v''^u)^b \\
&= S
\end{aligned}$$

This establishes that the unmodified SRP service algorithm, as well as the unmodified networking protocol, can work with the modified client that this formalism introduces. The one thing to care for is that all clients adhere to the same new formalism, since they will need to base their work on v'' and not v or v' . Since the password P is concealed by PKCS #11, this constraint is already implied by the technology.

5 Implementation as SRP #11

The following describes how to implement the last iteration of the SRP mechanism and demonstrates its integration with PKCS #11. Where the distinction is useful, we will distinguish ‘plain SRP’ from this modified form, which we will refer to as SRP #11.

Password establishment. A new ‘password’ is generated with the CKM_DH_PKCS_KEY_PAIR_GEN mechanism in `C_GenerateKeyPair()`, leading to a randomly generated private key P and a public key p . Only for reasons of backup might wrapped export of p be permissible, but P can be handled with more relaxation. The two are never used separately, but the public key may be used outside of PKCS #11 for reasons of efficiency.

The base value used is the generator g , to establish the relationship $p = g^P$. Both the base value and the modulus are commonly used group parameters for the SRP scheme, with a few standardised in Appendix A of RFC 5054. These must be supplied as parameters to `C_GenerateKeyPair()`. For standard modular-exponentiation Diffie-Hellman key exchange, cryptoanalytic advances suggest the use of non-standard groups, to evade dangers caused by pre-computed lookup tables.

Private modular exponentiation. The private key can be used as an exponent in modular exponentiation through `C_DeriveKey()` with the CKM_DH_PKCS_DERIVE method. The outcome of this operation can be stored in a CKO_SECRET_KEY of type CKK_GENERIC_SECRET with CKA_VALUE_LEN set to the number of bytes in the modulus. The template used must set CKA_EXTRACTABLE

to CK_TRUE, so the outcome can be retrieved from the session; the outcome is stored as a session key and after extraction this key should be deleted with C_DestroyObject() if the session persists. Programmers must be aware that implementations can remote leading zero bytes, and thus shorten the exported key.

Public key regeneration. It is possible to regenerate the public key p as g^P because the private key is annotated with non-sensitive CKA_PRIME holding the modulus and CKA_BASE holding g . The regenerated public key could be a session key that is not stored on the token even if the private key is. This means that it is possible to compute the public key from scratch, had it been deleted, and implement caching schemes and so on for it. This method can be used to save PKCS #11 storage capacity, or to avoid having to keep track of both keys. It is costly however; it takes an extra modular exponentiation operation. Another possibility is to export the public key value and store it off-token.

Credentials generation. When fresh Diffie-Hellman group parameters are desired, these can be generated first. To derive service access credentials from the password, generate a random salt s of the same size as the outcome of $H()$. The value of s is public, and may be generated on or off the PKCS #11 token, although the entropy of C_GenerateRandom() is likely to be of higher quality than more or less deterministic source from an operating system. Then, a hash function (not necessarily using the token's C_DigestXXX() functions) must be used to compute the hash value $x = H(s)$; the outcome can be used to calculate $p^{H(s)}$ off-token. The result of that operation is input to private modular exponentiation, which requires access to the token, to generate the outcome v'' . This outcome is supplied in a pair $\langle s, v \rangle$ to the service, together with the client's username C that serves as a lookup key for the pair.

Authentication. The client starts off generating random bytes, at least the size of the hash $H()$, and calls it a . Random bytes are best generated with C_GenerateRandom(), because the API potentially speaks to cryptographic hardware, making it yield at least as good random material as an operating system can. The client then uses private modular exponentiation to map a to $A = g^a$, which is transmitted to the service.

The service computes $k = H(N, g)$ and generates $B = k \cdot v'' + g^b$ for a local random secret b and sends B to the client. Note that the computation of B is completely done in modulo arithmetic. We are assuming SRP 6a, which does not generate a random value u on the server and sends it to the client, but instead both compute $u = H(A, B)$. Leading zero bytes should not be trimmed, to make A and B always have the same byte count as (the shortest representation of) the modulus.

Through the values of p and P , the client can now reconstruct v'' , which requires the non-extractable secret P ; moreover, the outcome cannot be reversed to find the private value P due to the properties of the discrete logarithm problem. The value v'' is also known to the service, so it may be exported from the PKCS #11 interface. The client can compute $k = H(N, g)$ and now uses v'' find $B - k \cdot v''$ with modular arithmetic, which is entirely based on values that the service also sees. The outcome equals g^b , which can be combined with session key a in `C.DeriveKey()` to find g^{ab} , a normal DH result.

The remainder of the modular computation of S'' consists of using the value $B - k \cdot v''$ in off-token multiplication and modular exponentiation, all based on the public values $B - k \cdot v''$, u , $H(s)$ and the locally generated random value a , all of which are safe to reveal outside of PKCS #11, and there will be two modular exponentiations with secret key P as the exponent, to be conducted on the token behind the protective PKCS #11 API. It is this last point where presence of the token is a strict requirement to derive the proper values.

Note that the value S'' is also available on the service, so we do not consider it a secret to be concealed behind PKCS #11. [**TODO:** How about the \dots^P and \dots^{PP} values, are they ‘public’ too?] Likewise, the session key $K = H(S)$ is also available to the service and not considered secret.

Note that the value of x'' is absent from the calculations. Instead, the secret key P is used. In a normal calculation the dependency on a secret key might be a liability, but when using PKCS #11 this is not as problematic because of its ability to encapsulate private keys. The value of x'' cannot be derived either, since the protected storage of the password values make it impossible to find the outcome of $H(s)p$ because that is not part of the DH mechanism.

Keys and Verifiers. The computation of $K'' = H(S'')$ is a trivial hashing operation that does not require padding of S'' , so we don’t do it in SRP #11. The computation of M_1'' and M_2'' are however sensitive to the problems of concatenation of variable-sized fields; bytes from one might be cut off and attached to another while retaining the same hash. Under SRP, this has not been solved with unambiguous separators, nor by length annotation, so the number of bytes of A and B in the hashing computations must be fixed to the same number of bytes as (the shortest representation of) the modulus. Sizes of K'' and M_1'' are fixed by the hash algorithm $H()$, so they need no special treatment. Finally, the value $H(N) \oplus H(g)$ does not have this problem and so there is no need to pad g ; both N and g can have the shortest possible byte sizes. It is common for SRP to apply no padding in the computation of $k = H(N, g)$ either.

6 Performance Considerations

The use of the more protective private-key encapsulation of PKCS #11 comes at a price of reduced efficiency. Note that this plays particularly at the client side, which means that it is subject to user selection of technology, and more importantly, meaning that scaling problems are not likely to be a problem — as we are talking of an alternative to manual password entry.

Still, a performance indication is useful. The overruling consideration in this sense is how many modular-exponential operations take place. More accurately put, the number of bits that may hold a value in the exponent of these operations. We are going to base our performance considerations on the assumption that the random values s , a , b , u all match the size of the outcome of the hash algorithm H . So, we can count the number of times such bits occur in the exponent of a calculation.

The chief concern is the authentication process. In the original process we find client-side exponents a , x and $a+ux$, sized like 1, 1 and 2 hashes, so plain SRP weighs like 4 modular exponentiations with the size of the outcome of H . In SRP #11, the client-side exponents during authentication are $H(s)$, a , a , u , P , P and $H(s)$, which weighs like 7 modular exponentiations.

Clearly, the inability of computing $a + ux$ before the exponentiation forces in extra computations in return for the added security of not handling x in SRP #11. The computational complexity of SRP #11 is about 75% higher than for plain SRP. This factor is probably irrelevant in comparison to the (possible) move of the calculation to a piece of hardware. Hardware tends to be either much slower (USB token) or much faster (HSM, crypto co-processor) at performing modular exponentiation than a CPU.

We performed measurements that compared an SRP implementation to SRP #11, the latter communicating with a software-based HSM. Both procedures included a client and a server running in the same process on the same machine, so as to eliminate network delays. We used the standard group with a 4095 bits modulus. The results indicate that SRP #11 takes ten to twelve times as much computation time as SRP. This has not been unravelled to separate the PKCS #11 overhead from the computations themselves, because there will always be PKCS #11 overhead when SRP #11 is being used. The total computation time of one SRP #11 verification run on a 2.2 GHz amd64 machine took about 175 ms.

7 Security Considerations

The values that are available as intermediate results in SRP #11 are all ‘public’, in the sense that the service can also find them based on what it knows. Most computations are based on a generator raised to a (secret or public) exponent, which is considered irreversable on grounds of the infeasibility of the discrete logarithm problem. [TODO: The values \dots^P and \dots^{P^P} might be non-‘public’; to be determined whether they could do harm.]

The value v'' is known to the service, and therefore its derivation as $p^{H(s)P}$ during authentication does not reveal information that should be protected by PKCS #11. The value $B - k \cdot v''$ is ‘public’, and so is a , so computing $(B - k \cdot v'')^a$ presents no problems. The same holds for $(B - k \cdot v'')^u$ and the next step, $((B - k \cdot v'')^u)^{H(s)}$. The latter value is then twice raised to the value P that is protected behind PKCS #11, but an attacker would need to crack discrete logarithms to reverse this to find P . The outcome of this last operation is equal to the service-side value S multiplied by the inverse of the public-value based $(B - k \cdot v'')^a$, so having the value once more reveals no new information. The remainder computers K , M_1 and M_2 on both ends, so there is no question about their revelation of new information that ought to be protected PKCS #11.

It is worth noting that a client might turn rogue when it is cracked; in such cases, it might gain the rights to authenticate, but it should not gain access to the passwords through which it can mount future attacks on other locations. This makes the common assumption that the credential protection of PKCS #11 is not cracked; an assumption that can be influenced by a suitable choice of PKCS #11 implementation. Plain SRP and SRP #11 are not different in this respect, and in fact it is common to all DH mechanism. Another point common to all three variations is that thrashing of the random values a and b suffices to thwart future attackers from recovering session keys of traffic preceding the crack on the client (forward secrecy).

The value $p = g^P$ is retained as a ‘cache’ value. Rogue access to PKCS #11 might consider replacing this value with a more suitable value, to influence future computations based on it. This is not likely to lead to anything useful however; first, it would be noticed when accessing a non-rogue service; second, the value P is still used in computations; and third, the discrete logarithm problem makes it notoriously difficult to derive a value of p that would lead to a desired outcome. A much better attack would be to replace both p and P which is comparable to replacing a password; this however, would not leak any information about the prior value of P and would invalidate access to services setup with it.

Modular exponentiation cannot be reversed, as a result of the discrete log-

arithm problem. The start of the computation is $B - k \cdot v''$, which equals g^b and that uses an unknown value b . Raising this to the powers u , $H(s)$, P and P in any order should stir the value even further, but in lieu of the value b on the client and due to the discrete logarithm problem it is unlikely that an attack could be mounted from this angle. **Open Question:** Can we prove that?

In conclusion, the maximum protection that PKCS #11 can deliver has been achieved with SRP #11, under the assumption that what the service knows is sufficiently public. That angle however, is subject of potential criticism. It could be argued that the service is in a special position. For one, it generates random string b internally and only publishes the corresponding B ; but that alone would not make much difference, an active MITM attack could produce similar values. The vital distinction is that the service knows v' , which is assumed to have been relayed to it in a covert manner, and so it is a shared secret between client and service — one that would not be available to anyone on the channel between client and service.

Unfortunately, it is not possible to redesign SRP #11 to treat v'' as a secret value. This can be seen directly in the computations; the construction $(B - k \cdot v'')^e$, regardless of the exponent e cannot be computed with DH values, because subtraction is not foreseen in this mechanism. This is where PKCS #11 reveals that it was not designed with plain SRP or SRP #11 in mind. And since exponentiation does not distribute through subtraction, neither modular or otherwise, we cannot construct any such mechanism cleverly.

Note that the visibility of v'' is less than perfect, but not truly problematic in itself. The sole reason for treating all data available to the service as public is to avoid putting too much faith in what PKCS #11 can do. Apparently, this is its limitation. Note that this is also caused by something else — the client is the party that originally generated v'' and would be able to do that again.

In conclusion, SRP #11 actually improves the security of SRP by taking the passwords out of sight and placing them behind PKCS #11; in addition, the password-equivalent value x is not available to the SRP #11 computations due to PKCS #11 shielding. The one thing that is not possible, is to also conceal the computation of verifier values v'' from the client, but it has been demonstrated that this would have introduced a false sense of security anyway. In conclusion, SRP #11 establishes precisely the maximum achievable and the maximum desirable protection of SRP credentials.

8 Conclusions

This article introduces SRP #11, a variation on SRP that uses PKCS #11 to protect the password and password-equivalent values that are computed on the client side. It has been shown that a direct implementation of plain SRP with PKCS #11 does protect the original password, but not the password-equivalent value x , which means it is useless. In contrast, the SRP #11 variation achieves precisely the maximum achievable and maximum desirable that PKCS #11 could offer.

It has been shown that the modified generation of salt and verifier combined with the modified client-side authentication computations compute values equivalent to those of plain SRP, and that the service or protocol require no changes. In other words, only the client needs to modify its computations if it intends to use PKCS #11.

In terms of efficiency, the SRP #11 mechanism adds a factor ten to twelve to the running time in comparison to plain SRP, but the execution times are still quite acceptable. When PKCS #11 is implemented on hardware, then choice is likely to have a stronger impact on the speed of the computations — they would take longer when a plain USB token is used, and would go down when cryptographic acceleration is built into the hardware.